

test

Huanyu Chen

2024-04-22

```
# Load necessary libraries
```

```
library(mclust)
```

```
## Warning: package 'mclust' was built under R version 4.3.2
```

```
## Package 'mclust' version 6.1
```

```
## Type 'citation("mclust")' for citing this R package in publications.
```

```
library(tidyverse)
```

```
## Warning: package 'tidyr' was built under R version 4.3.2
```

```
## -- Attaching core tidyverse packages ----- tidyverse 2.0.0 --
```

```
## v dplyr      1.1.4      v readr      2.1.5
```

```
## v forcats   1.0.0      v stringr   1.5.1
```

```
## v ggplot2    3.4.4      v tibble    3.2.1
```

```
## v lubridate  1.9.3      v tidyr     1.3.1
```

```
## v purrr      1.0.2
```

```
## -- Conflicts ----- tidyverse_conflicts() --
```

```
## x dplyr::filter() masks stats::filter()
```

```
## x dplyr::lag()     masks stats::lag()
```

```
## x purrr::map()     masks mclust::map()
```

```
## i Use the conflicted package (http://conflicted.r-lib.org/) to force all conflicts to become errors
```

```
# Read the data
```

```
df <- read.csv("./data/diabetes.csv") |>
```

```
  janitor::clean_names() |>
```

```
  mutate(across(c(glucose, blood_pressure, skin_thickness, insulin, bmi), ~na_if(.x, 0)))
```

```
# Handle missing values
```

```
df[is.na(df)] <- 0
```

```
# Extract features
```

```
X <- df[, c("pregnancies", "glucose", "blood_pressure", "skin_thickness", "insulin", "bmi",  
           "diabetes_pedigree_function", "age")]
```

```
# Define the EM algorithm function
```

```
EM_algorithm <- function(X, num_clusters, max_iter = 100, tol = 1e-6) {
```

```

# Initialize parameters
model <- Mclust(X, G = num_clusters)

for (iter in 1:max_iter) {
  # E-step
  responsibilities <- matrix(0, nrow = nrow(X), ncol = num_clusters)
  for (k in 1:num_clusters) {
    responsibilities[, k] <- model$z[, k] * model$parameters$pro[k]
  }
  responsibilities <- responsibilities / rowSums(responsibilities)

  # M-step
  model <- Mclust(X, G = num_clusters, z = responsibilities)

  # Calculate log-likelihood and check for convergence
  log_likelihood <- sum(log(apply(model$z * model$parameters$pro, 1, sum)))
  if (iter > 1 && abs(log_likelihood - prev_log_likelihood) < tol) {
    break
  }
  prev_log_likelihood <- log_likelihood
}

return(model)
}

# Call the EM algorithm
num_clusters <- 3 # Set the number of clusters
gmm_model <- EM_algorithm(X, num_clusters)

# Get the cluster assignments for each sample
clusters <- predict(gmm_model)

# View the feature patterns of the clusters
cluster_means <- t(apply(gmm_model$parameters$mean, 2, function(x) round(x, 2)))
print(cluster_means)

```

```

##      pregnancies glucose blood_pressure skin_thickness insulin   bmi
## [1,]         4.25  133.59           62.55          25.76  137.72  32.63
## [2,]         3.11  112.17           69.86          28.18   88.94  31.97
## [3,]         4.86  124.39           74.63           0.00    0.00  31.35
##      diabetes_pedigree_function   age
## [1,]                   0.68 38.07
## [2,]                   0.40 28.05
## [3,]                   0.40 38.23

```

```

# View the mixing weights of the clusters
mixing_weights <- round(gmm_model$parameters$pro, 2)
print(mixing_weights)

```

```

## [1] 0.27 0.49 0.25

```

```
# Further analysis can be performed based on the cluster assignments,
# for example, analyzing the relationship between cluster assignments and diabetes status
```

```
library(ggplot2)
```

```
# Ensure the cluster and diabetes_status variables are factors
```

```
df$cluster <- factor(clusters$classification)
```

```
cluster_means <- aggregate(outcome ~ cluster, data = df, FUN = mean)
```

```
# Plot the distribution of diabetes status within each cluster
```

```
ggplot(df, aes(x = cluster, fill = factor(outcome))) +
```


```
  geom_bar(position = "fill") +
```

```
  geom_text(data = cluster_means, aes(label = sprintf("%.2f", outcome), y = outcome), vjust = -0.5, col
```

```
  scale_fill_manual(values = c("0" = "blue", "1" = "red"), labels = c("0" = "Non-Diabetes", "1" = "Diab
```

```
  labs(x = "Cluster", y = "Proportion", fill = "Diabetes Status") +
```

```
  ggtitle("Distribution of Diabetes Status by Cluster (Proportion by Mean Outcome)")
```



```
# OR
```

```
# Plot the mean outcome for each cluster
```


```
ggplot(cluster_means, aes(x = cluster, y = outcome)) +
```

```
  geom_bar(stat = "identity", fill = "skyblue") +
```

```
  geom_text(aes(label = sprintf("%.2f", outcome)), vjust = -0.5, color = "black", size = 3) +
```

```
  labs(x = "Cluster", y = "Mean Outcome") +
```

```
  ggtitle("Mean Outcome for Each Cluster")
```



```
# Plot the relationship between glucose and insulin, colored by cluster  
ggplot(df, aes(x = glucose, y = insulin, color = factor(cluster))) +  
  geom_point() +  
  labs(x = "Glucose", y = "Insulin", color = "Cluster") +  
  ggtitle("Relationship between Glucose and Insulin by Cluster")
```

Relationship between Glucose and Insulin by Cluster

