

# test

Huanyu Chen

2024-04-22

```
# Load necessary libraries
```

```
library(mclust)
```

```
## Warning: package 'mclust' was built under R version 4.3.2
```

```
## Package 'mclust' version 6.1
```

```
## Type 'citation("mclust")' for citing this R package in publications.
```

```
library(tidyverse)
```

```
## Warning: package 'tidyr' was built under R version 4.3.2
```

```
## -- Attaching core tidyverse packages ----- tidyverse 2.0.0 --
```

```
## v dplyr      1.1.4      v readr      2.1.5
```

```
## v forcats    1.0.0      v stringr    1.5.1
```

```
## v ggplot2    3.4.4      v tibble     3.2.1
```

```
## v lubridate  1.9.3      v tidyr      1.3.1
```

```
## v purrr      1.0.2
```

```
## -- Conflicts ----- tidyverse_conflicts() --
```

```
## x dplyr::filter() masks stats::filter()
```

```
## x dplyr::lag()     masks stats::lag()
```

```
## x purrr::map()     masks mclust::map()
```

```
## i Use the conflicted package (http://conflicted.r-lib.org/) to force all conflicts to become errors
```

```
# Read the data
```

```
df <- read.csv("./data/diabetes.csv") |>
```

```
  janitor::clean_names() |>
```

```
  mutate(across(c(glucose, blood_pressure, skin_thickness, insulin, bmi), ~na_if(.x, 0)))
```

```
# Handle missing values
```

```
df[is.na(df)] <- 0
```

```
# Extract features
```

```
X <- df[, c("pregnancies", "glucose", "blood_pressure", "skin_thickness", "insulin", "bmi",  
           "diabetes_pedigree_function", "age")]
```

```
# Define the EM algorithm function
```

```
EM_algorithm <- function(X, num_clusters, max_iter = 100, tol = 1e-6) {
```

```

# Initialize parameters
model <- Mclust(X, G = num_clusters)

# Initialize variables to track convergence
prev_log_likelihood <- -Inf

for (iter in 1:max_iter) {
  # E-step
  responsibilities <- matrix(0, nrow = nrow(X), ncol = num_clusters)
  for (k in 1:num_clusters) {
    responsibilities[, k] <- model$z[, k] * model$parameters$pro[k]
  }
  responsibilities <- responsibilities / rowSums(responsibilities)

  # M-step
  model <- Mclust(X, G = num_clusters, z = responsibilities)

  # Calculate log-likelihood
  log_likelihood <- sum(log(apply(model$z * model$parameters$pro, 1, sum)))

  # Check for convergence
  if (abs(log_likelihood - prev_log_likelihood) < tol) {
    break
  }

  # Update previous log-likelihood
  prev_log_likelihood <- log_likelihood
}

return(model)
}

# Call the EM algorithm
num_clusters <- 2 # Set the number of clusters
gmm_model <- EM_algorithm(X, num_clusters)

# Get the cluster assignments for each sample
clusters <- predict(gmm_model)

# View the feature patterns of the clusters
cluster_means <- t(apply(gmm_model$parameters$mean, 2, function(x) round(x, 2)))
print(cluster_means)

##      pregnancies glucose blood_pressure skin_thickness insulin   bmi
## [1,]         5.31  130.41          69.31          19.32   94.14  32.99
## [2,]         2.06  109.29          68.86          22.02   62.32  30.77
##      diabetes_pedigree_function   age
## [1,]                0.55  39.64
## [2,]                0.38  25.45

# View the mixing weights of the clusters
mixing_weights <- round(gmm_model$parameters$pro, 2)
print(mixing_weights)


```

```
## [1] 0.55 0.45
```

```
# Further analysis can be performed based on the cluster assignments,  
# for example, analyzing the relationship between cluster assignments and diabetes status
```

```
# Define a function to run EM algorithm with different numbers of clusters  
run_EM_with_different_clusters <- function(X, max_clusters = 10, max_iter = 100, tol = 1e-6) {  
  likelihoods <- numeric(max_clusters)  
  
  for (num_clusters in 2:max_clusters) {  
    # Run EM algorithm  
    model <- EM_algorithm(X, num_clusters, max_iter, tol)  
  
    # Calculate log-likelihood  
    log_likelihood <- sum(log(apply(model$z * model$parameters$pro, 1, sum)))  
  
    # Store log-likelihood  
    likelihoods[num_clusters] <- log_likelihood  
  }  
  
  return(likelihoods)  
}  
  
# Call the function to run EM with different numbers of clusters  
likelihoods <- run_EM_with_different_clusters(X)  
  
# Plot the likelihoods for different numbers of clusters  
plot(1:length(likelihoods), likelihoods, type = "b", xlab = "Number of Clusters", ylab = "Log-Likelihood")
```

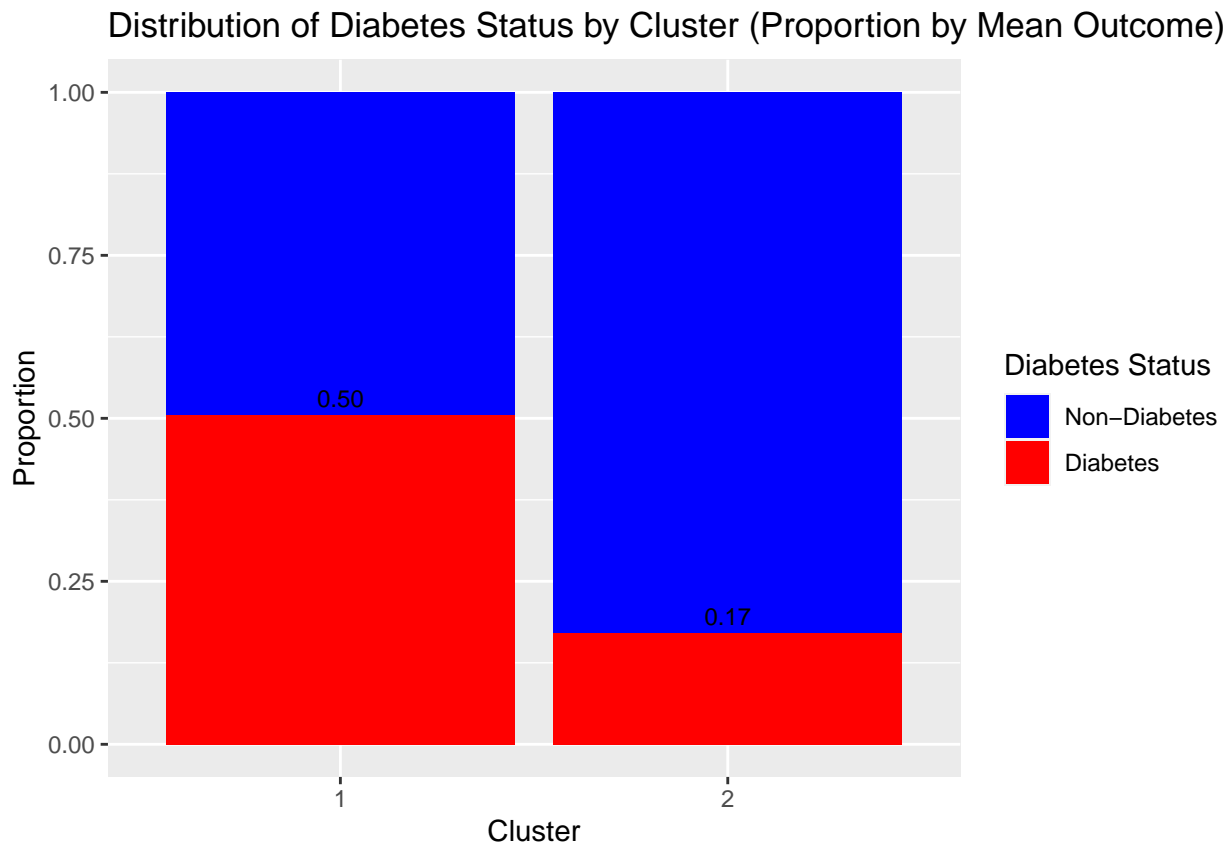
## Log-Likelihood vs. Number of Clusters



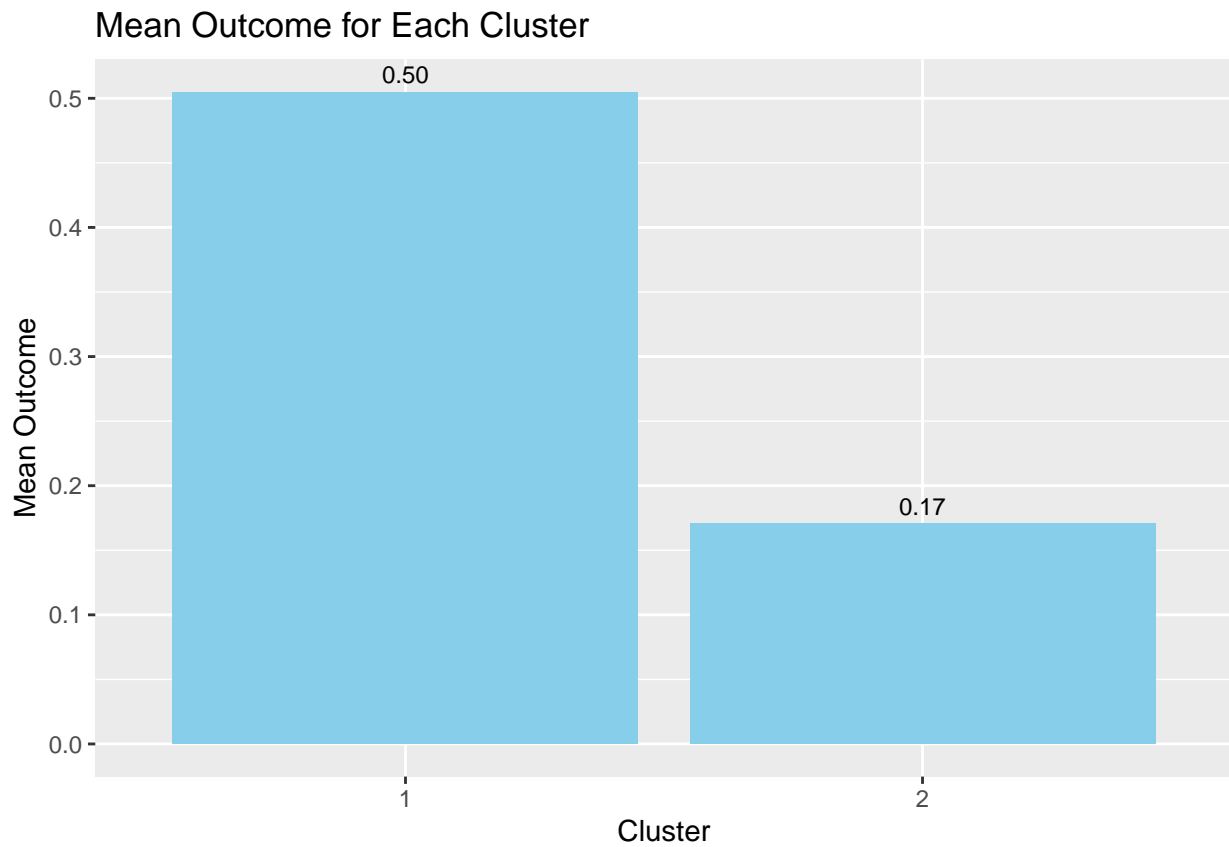
```
library(ggplot2)

# Ensure the cluster and diabetes_status variables are factors
df$cluster <- factor(clusters$classification)
cluster_means <- aggregate(outcome ~ cluster, data = df, FUN = mean)

# Plot the distribution of diabetes status within each cluster
ggplot(df, aes(x = cluster, fill = factor(outcome))) +
  geom_bar(position = "fill") +
  geom_text(data = cluster_means, aes(label = sprintf("%.2f", outcome), y = outcome), vjust = -0.5, color = "black", size = 3) +
  scale_fill_manual(values = c("0" = "blue", "1" = "red"), labels = c("0" = "Non-Diabetes", "1" = "Diabetes")) +
  labs(x = "Cluster", y = "Proportion", fill = "Diabetes Status") +
  ggtitle("Distribution of Diabetes Status by Cluster (Proportion by Mean Outcome)")
```



```
# OR
# Plot the mean outcome for each cluster
ggplot(cluster_means, aes(x = cluster, y = outcome)) +
  geom_bar(stat = "identity", fill = "skyblue") +
  geom_text(aes(label = sprintf("%.2f", outcome)), vjust = -0.5, color = "black", size = 3) +
  labs(x = "Cluster", y = "Mean Outcome") +
  ggtitle("Mean Outcome for Each Cluster")
```



```
# Plot the relationship between glucose and insulin, colored by cluster
ggplot(df, aes(x = glucose, y = insulin, color = factor(cluster))) +
  geom_point() +
  labs(x = "Glucose", y = "Insulin", color = "Cluster") +
  ggtitle("Relationship between Glucose and Insulin by Cluster")
```

